

B. Testing Notebook

After training the model at [train.ipynb](#) we need to test its accuracy

1) *Imports*: We begin by importing our dependencies, for testing, we need tensorflow and tensorflow_datasets much like we did with training, but we also need numpy and seaborn

```
[1]: import tensorflow as tf
import numpy as np
import seaborn as sns
from pprint import pprint
import matplotlib.pyplot as plt
import tensorflow_datasets as tfds
```

2) *The dataset*: We use the `load` method to load the mnist dataset, but this time, we only load the test split which contains 10000 images

```
[ ]: # load the mnist dataset
dstest, dsinfo = tfds.load(
    'mnist',
    split=['test'], # only need the test set
    data_dir='../dataset/',
    shuffle_files=True,
    as_supervised=True,
    with_info=True,
)

dstest = dstest[0] # Because tfds.load returns a list
```

3) *Preprocessing*: We follow the same preprocessing we did while training the model, which is practically nothing that is necessary. However for performance reasons, we autotune and batch the test set

```
[3]: batch_size = 128

# Evaluation pipeline
dstest = dstest.batch(batch_size)
dstest = dstest.cache()
dstest = dstest.prefetch(tf.data.AUTOTUNE)
```

4) *The model*: We load the model created in the train script

```
[4]: model = tf.keras.models.load_model('./model.h5')
```

Printing the class names

```
[5]: class_names = dsinfo.features['label'].names
print(class_names)
```

```
['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
```

Generating the model predictions

```
[6]: model_probabilities = model.predict(dstest)
pprint(model_probabilities)
```

```
2022-04-24 23:30:08.887871: I tensorflow/stream_executor/cuda/cuda_dnn.cc:368]
Loaded cuDNN version 8303
```

```
array([[1.3535260e-12, 9.2269102e-11, 1.0000000e+00, ..., 1.2002113e-10,
        3.6733475e-08, 1.2856706e-10],
```

```
[9.9999809e-01, 1.9031241e-09, 3.8650558e-08, ..., 4.1776258e-07,
 1.7634169e-07, 8.7255484e-07],
[1.0604630e-11, 4.1063618e-08, 1.8023245e-09, ..., 8.3668425e-07,
 8.0477860e-11, 7.3748092e-09],
...,
[1.1503052e-11, 2.6557379e-10, 5.7338542e-07, ..., 1.2051790e-10,
 9.9999917e-01, 2.3892071e-07],
[9.9994540e-01, 1.3106036e-07, 2.7527710e-08, ..., 6.4490371e-07,
 3.2164176e-09, 5.1208815e-05],
[1.8934915e-10, 3.5097508e-10, 2.8572342e-11, ..., 4.0571066e-12,
 3.7729077e-08, 2.4009706e-08]], dtype=float32)
```

`model.predict` returns a list of lists, where each inner list contains the probabilities that image belongs to each class, to get the predicted labels we choose the class with the highest probability

```
[7]: predictions = [np.argmax(x) for x in model_probabilities]
      pprint(predictions[:10]) # printing the last 10 results
```


```
[2, 0, 4, 8, 7, 6, 0, 6, 3, 1]
```

5) *Evaluation of the model:* We get the actual true labels using this one line

```
[8]: labels = np.concatenate([y for x, y in dtest], axis=0)
```

We then generate a confusion matrix and plot it

```
[9]: confusion_matrix = tf.math.confusion_matrix(
      labels=labels,
      predictions=predictions,
      )
      sns.heatmap(confusion_matrix,
                  annot=True,
                  xticklabels=class_names,
                  yticklabels=class_names,
                  fmt='g')
      plt.xlabel('Predicted')
      plt.ylabel('True')
      plt.savefig('../paper/figs/confusion_matrix.svg', format='svg')
      plt.show()
```



With the confusion matrix we can calculate the accuracy as

$$\frac{\sum_{i=0}^9 k_{ii}}{\sum_{i=0}^9 \sum_{j=0}^9 k_{ij}} \times 100$$

Where k_{xy} represent an element in the x^{th} row and the y^{th} column in the confusion matrix, in other words it is the sum of the elements in the diagonal of the confusion matrix, divided by the total sum.

```
[10]: diagonal_sum = 0
total_sum = 0

for i in range(len(class_names)):
    for j in range(len(class_names)):
        if(i == j):
            diagonal_sum += confusion_matrix[i][j]
            total_sum += confusion_matrix[i][j]

print("Diagonal sum: {}, Total sum: {}".format(diagonal_sum, total_sum)) #_
    ↳ total sum should be 10000 as the test split is 10000 images
accuracy = 100 * diagonal_sum/total_sum
print("Accuracy: {}".format(accuracy))
```

Diagonal sum: 9908, Total sum: 10000
Accuracy: 99.08

We can also calculate the per-digit accuracy for a digit i as

$$\frac{k_{ii}}{\sum_{j=0}^9 k_{ij}} \times 100$$

(For example, for the digit 0, it was predicted correctly 970+0+0+0+0+0+5+2+1+2) With that information, we can plot a bar chart to see how accurately our model predicts each digit with the following code

```
[12]: fig = plt.figure(figsize=(15, 5))
ax = fig.add_axes([0, 0, 0.7, 1])
digits = list(range(0, 10))
digit_acc = []
for i in digits:
    row_sum = sum([confusion_matrix[i][j] for j in digits])
    digit_acc.append(100 * confusion_matrix[i][i]/row_sum)
plt.yticks(digits)
bars = ax.barh(digits, digit_acc)
ax.bar_label(bars)
ax.spines["top"].set_bounds(0, 113)
ax.spines["bottom"].set_bounds(0, 113)
ax.spines["right"].set_position(("outward", 60))
plt.rcParams.update({'font.size': 22}) # Increasing the text size
fig.tight_layout()
plt.savefig('../paper/figs/bar_plot.svg', format='svg', bbox_inches='tight')
plt.show()
```

/tmp/ipykernel_24732/3498705007.py:15: UserWarning: This figure includes Axes that are not compatible with tight_layout, so results might be incorrect.

```
fig.tight_layout()
```

