

APPENDIX

A. Training Notebook

This notebook contains the code for training the model on the MNIST dataset, evaluation can be found at [test.ipynb](#)

1) *Imports*: For training the model, we need only tensorflow and tensorflow_datasets, which will be used to retrieve the MNIST dataset

```
[1]: import matplotlib.pyplot as plt
import tensorflow as tf
import tensorflow_datasets as tfds
```

2) *Loading the dataset*: We use the `load` method to load the mnist dataset

```
[ ]: # Import the dataset mnist is 60k images of 28x28 pixels
# And 10k images for testing
(dstrain, dtest), dsinfo = tfds.load(
    'mnist',
    split=['train', 'test'],
    data_dir='../dataset/',
    shuffle_files=True,
    as_supervised=True,
    with_info=True,
)
```

Print the image shape and class names

```
[3]: # Summarize loaded datasets
print('\nDataset info:')
print('Image shape:')
print(dsinfo.features['image'].shape)
print('Class Names')
print(dsinfo.features['label'].names)
```

Dataset info:

Image shape:

(28, 28, 1)


Class Names

['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']

Visualize an image from the dataset

```
[4]: # Visualize a single image
def visualize_image(image, label):
    plt.imshow(image, cmap='gray')
    plt.title('Label: {}'.format(label))
    plt.show()

# Now use that function
mnist_example = dstrain.take(1)
for sample in mnist_example:
    image, label = sample[0], sample[1]
    visualize_image(image, label)
    break
```



2022-04-24 22:37:42.374927: W tensorflow/core/kernels/data/cache_dataset_ops.cc:768] The calling iterator did not fully read the dataset being cached. In order to avoid unexpected truncation of the dataset, the partially cached contents of the dataset will be discarded. This can happen if you have an input pipeline similar to ``dataset.cache().take(k).repeat()``. You should use ``dataset.take(k).cache().repeat()`` instead.

3) *Preprocessing*: We preprocess the data by batching it, and applying an autotuned prefetch to make the fetching of the data faster

Note: instead of preprocessing the data right now, we have a resizing layer in the model architecture

From the Tensorflow Documentation: > Prefetching overlaps the preprocessing and model execution of a training step. While the model is executing training step s , the input pipeline is reading the data for step $s+1$. Doing so reduces the step time to the maximum (as opposed to the sum) of the training and the time it takes to extract the data.

```
[5]: batch_size = 128

dstrain = dstrain.batch(batch_size)
dstrain = dstrain.cache()
dstrain = dstrain.shuffle(dsinfo.splits['train'].num_examples)
dstrain = dstrain.prefetch(tf.data.AUTOTUNE)
```

We do the same for the testing set

```
[6]: dstest = dstest.batch(batch_size)
dstest = dstest.cache()
dstest = dstest.prefetch(tf.data.AUTOTUNE)
```

4) *The Model:* We then create the model architecture

```
[7]: model = tf.keras.Sequential([
    tf.keras.layers.Rescaling(1./255, input_shape=(28, 28, 1)),
    tf.keras.layers.Conv2D(28, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv2D(28, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

We compile the model, where we use the adam optimizer, which is a version of stochastic gradient descent, and cross categorical cross entropy for the loss function which is mathematically written as:

$$J(\mathbf{w}) = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

We then build the model, and print the summary

```
[8]: model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'],
)

# build the model
model.build()
# print the model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 28, 28, 1)	0
conv2d (Conv2D)	(None, 26, 26, 28)	280
max_pooling2d (MaxPooling2D)	(None, 13, 13, 28)	0
batch_normalization (Batch Normalization)	(None, 13, 13, 28)	112
conv2d_1 (Conv2D)	(None, 11, 11, 28)	7084

max_pooling2d_1 (MaxPooling 2D)	(None, 5, 5, 28)	0
batch_normalization_1 (Batch Normalization)	(None, 5, 5, 28)	112
flatten (Flatten)	(None, 700)	0
dense (Dense)	(None, 128)	89728
dense_1 (Dense)	(None, 128)	16512
dense_2 (Dense)	(None, 128)	16512
dense_3 (Dense)	(None, 10)	1290

=====

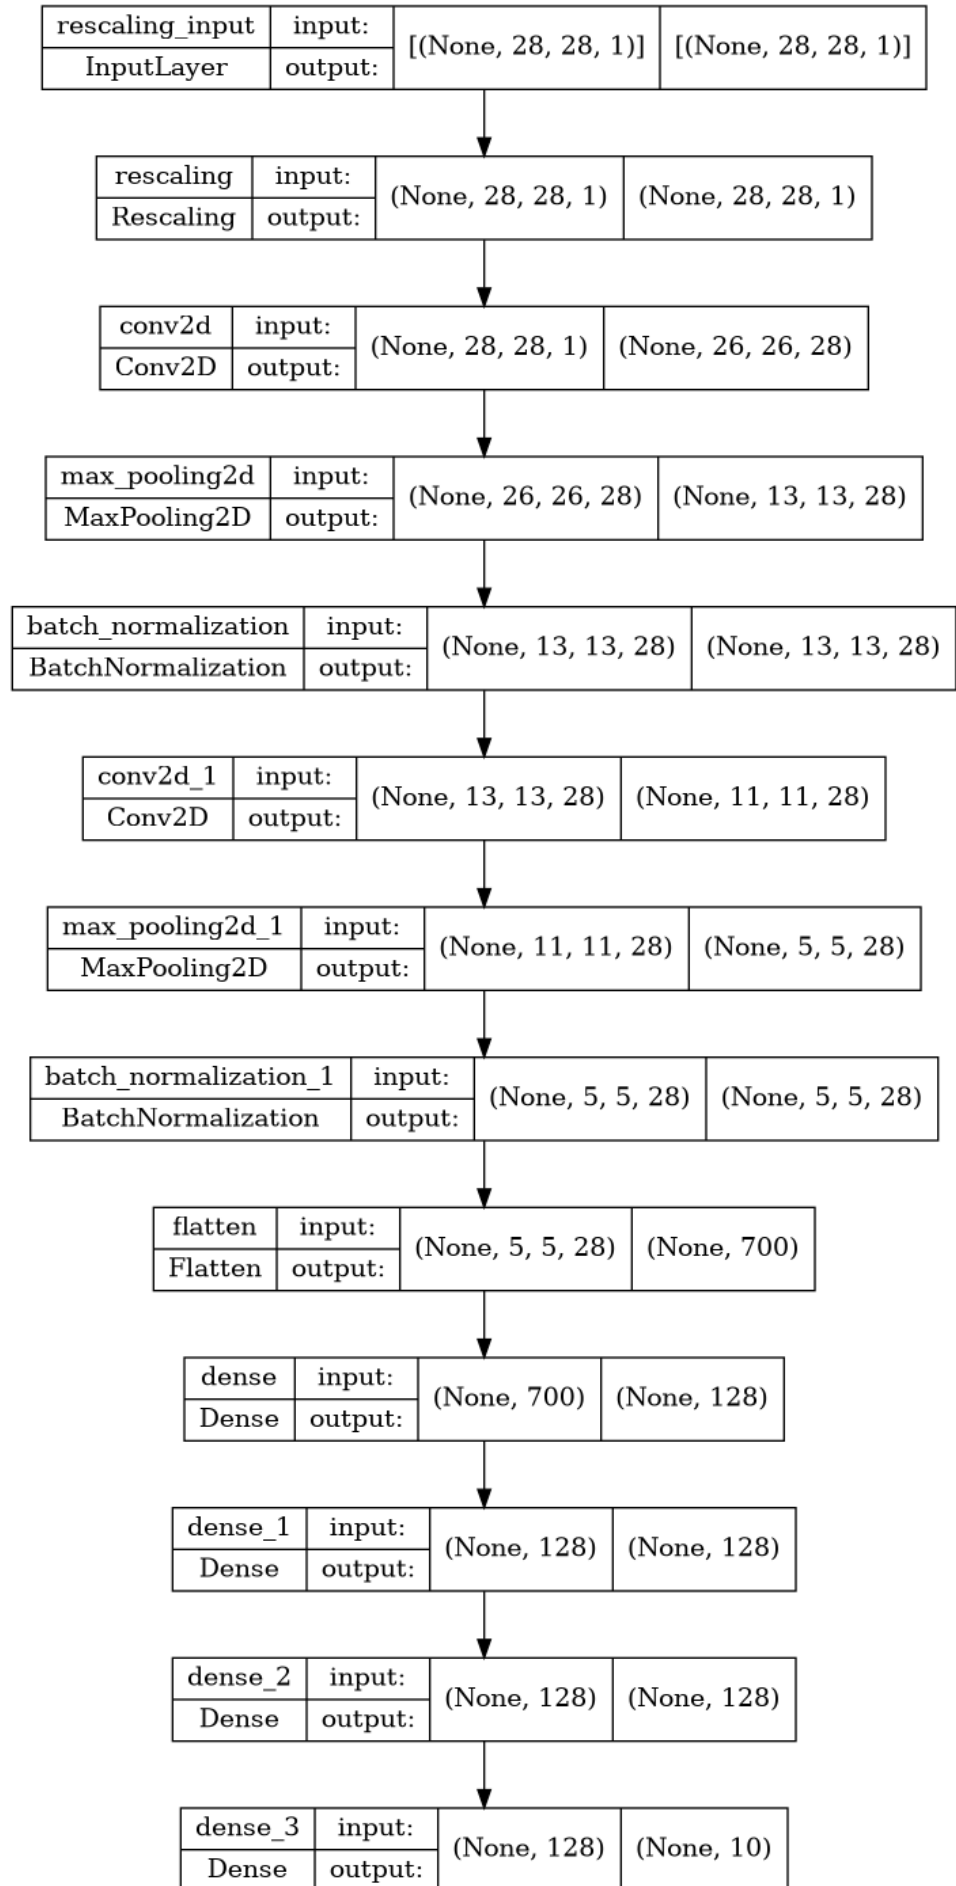
Total params: 131,630
Trainable params: 131,518
Non-trainable params: 112

Create the architecture diagram and save it to an svg

```
[9]: import pydot
from IPython.display import Image, display

def view_pydot(pdot):
    plt = Image(pdot.create_png())
    display(plt)

architecture = tf.keras.utils.model_to_dot(model, show_shapes = True)
view_pydot(architecture)
architecture.write_svg('../paper/figs/cnn_architecture.svg')
```



We then train the model, with an early stopping callback, and saving the best model along the way

```
[10]: history = model.fit(
    dstrain,
    epochs=30,
    validation_data=dstest,
    callbacks=[
        tf.keras.callbacks.EarlyStopping(
            monitor='val_loss',
            patience=5,
        ),
        tf.keras.callbacks.ModelCheckpoint(
            'model.h5',
            monitor='val_loss',
            save_best_only=True,
            verbose=1)
    ]
)
```

Epoch 1/30

2022-04-24 22:38:03.283133: I tensorflow/stream_executor/cuda/cuda_dnn.cc:368]
Loaded cuDNN version 8303

461/469 [=====>.] - ETA: 0s - loss: 0.1330 - accuracy: 0.9598

Epoch 1: val_loss improved from inf to 0.21989, saving model to model.h5

469/469 [=====] - 8s 5ms/step - loss: 0.1315 - accuracy: 0.9603 - val_loss: 0.2199 - val_accuracy: 0.9304

Epoch 2/30

465/469 [=====>.] - ETA: 0s - loss: 0.0392 - accuracy: 0.9878

Epoch 2: val_loss improved from 0.21989 to 0.03657, saving model to model.h5

469/469 [=====] - 2s 5ms/step - loss: 0.0391 - accuracy: 0.9878 - val_loss: 0.0366 - val_accuracy: 0.9877

Epoch 3/30

469/469 [=====] - ETA: 0s - loss: 0.0267 - accuracy: 0.9918

Epoch 3: val_loss improved from 0.03657 to 0.03342, saving model to model.h5

469/469 [=====] - 2s 5ms/step - loss: 0.0267 - accuracy: 0.9918 - val_loss: 0.0334 - val_accuracy: 0.9891

Epoch 4/30

466/469 [=====>.] - ETA: 0s - loss: 0.0196 - accuracy: 0.9939

Epoch 4: val_loss improved from 0.03342 to 0.03083, saving model to model.h5

469/469 [=====] - 2s 5ms/step - loss: 0.0195 - accuracy: 0.9939 - val_loss: 0.0308 - val_accuracy: 0.9908

Epoch 5/30

467/469 [=====>.] - ETA: 0s - loss: 0.0175 - accuracy: 0.9942

Epoch 5: val_loss did not improve from 0.03083

469/469 [=====] - 2s 5ms/step - loss: 0.0175 - accuracy: 0.9942 - val_loss: 0.0458 - val_accuracy: 0.9880

```

Epoch 6/30
466/469 [=====>.] - ETA: 0s - loss: 0.0143 - accuracy:
0.9959
Epoch 6: val_loss did not improve from 0.03083
469/469 [=====] - 2s 5ms/step - loss: 0.0143 -
accuracy: 0.9959 - val_loss: 0.0458 - val_accuracy: 0.9866
Epoch 7/30
466/469 [=====>.] - ETA: 0s - loss: 0.0124 - accuracy:
0.9958
Epoch 7: val_loss did not improve from 0.03083
469/469 [=====] - 2s 5ms/step - loss: 0.0124 -
accuracy: 0.9958 - val_loss: 0.0410 - val_accuracy: 0.9898
Epoch 8/30
465/469 [=====>.] - ETA: 0s - loss: 0.0105 - accuracy:
0.9967
Epoch 8: val_loss did not improve from 0.03083
469/469 [=====] - 2s 5ms/step - loss: 0.0106 -
accuracy: 0.9966 - val_loss: 0.0406 - val_accuracy: 0.9889
Epoch 9/30
467/469 [=====>.] - ETA: 0s - loss: 0.0100 - accuracy:
0.9968
Epoch 9: val_loss did not improve from 0.03083
469/469 [=====] - 2s 5ms/step - loss: 0.0100 -
accuracy: 0.9968 - val_loss: 0.0404 - val_accuracy: 0.9886

```

Plot the accuracies and the losses across the range of epochs

```

[11]: accuracy = history.history['accuracy']
      val_accuracy = history.history['val_accuracy']

      loss = history.history['loss']
      val_loss = history.history['val_loss']

      epochs_range = range(len(accracy))

      plt.figure(figsize=(4, 8))
      plt.plot(epochs_range, accuracy, label='Training Accuracy')
      plt.plot(epochs_range, val_accuracy, label='Validation Accuracy')
      plt.legend(loc='lower right')
      plt.savefig('../paper/figs/accuracy.svg', format='svg')
      plt.show()

      plt.figure(figsize=(4, 8))
      plt.plot(epochs_range, loss, label='Training Loss')
      plt.plot(epochs_range, val_loss, label='Validation Loss')
      plt.legend(loc='upper right')
      plt.savefig('../paper/figs/loss.svg', format='svg')
      plt.show()

```

